




Machine Learning and Data Mining

Lecture 3.1: Perceptron, Evaluating Classifiers, Naïve Bayes





RECAP


From Last Lecture



Beyond Lines but Still Linear?

Idea: Solve a linear regression problem in a feature space that is non-linear in the original input! For example, could add a x^2 term.

X		Y
	x	Y
1	-0.25	-4.47
1	0.9	11.08
1	0.46	-10.44
1	0.2	7.19
1	-0.69	7.08
1	-0.69	5.38
1	-0.88	11.55
1	0.73	10.85
1	0.2	7.85
1	0.42	-8.31



X'			Y
	x	x^2	Y
1	-0.25	0.06	-4.47
1	0.9	0.81	11.08
1	0.46	0.21	-10.44
1	0.2	0.04	7.19
1	-0.69	0.48	7.08
1	-0.69	0.48	5.38
1	-0.88	0.77	11.55
1	0.73	0.53	10.85
1	0.2	0.04	7.85
1	0.42	0.18	-8.31

Solve for w such that:

$$y_i = w_0 1 + w_1 x_i + w_2 x_i^2$$

Linear function of non-linear transformations of x

$$y_i = [1, x_i, x_i^2] \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix}$$

Let's generalize this further.

Each data point has d features. Let $\Phi: \mathbb{R}^d \rightarrow \mathbb{R}^{d'}$ be our basis function. Can make any arbitrary choice we want here based on how the data looks.

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}_{1 \times d}^T$$

$$\Phi_A(\mathbf{x}) = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ x_1 x_2 \\ \sin(x_1) \\ \cos(x_2) \end{bmatrix}^T$$

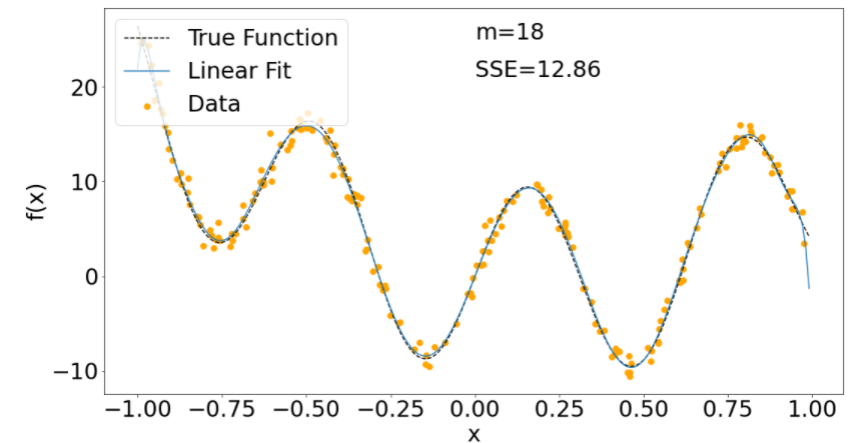
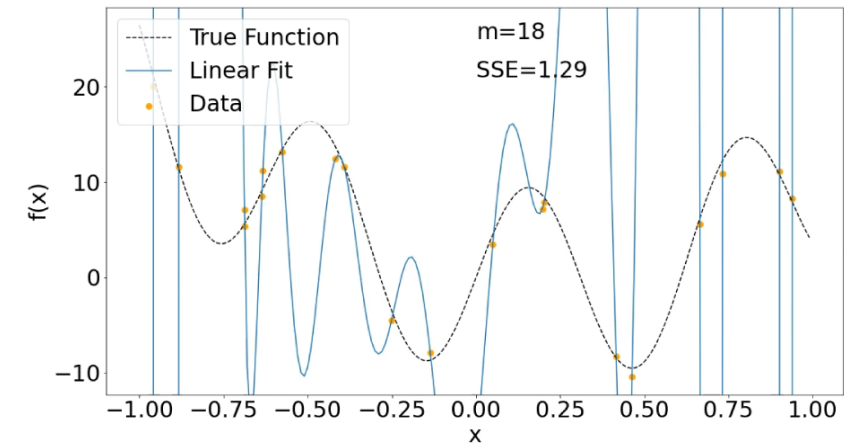
$$\Phi_B(\mathbf{x}) = \begin{bmatrix} 1 \\ x_1 x_2 \\ x_1 x_3 \\ x_1 x_4 \\ \vdots \\ x_n x_n \end{bmatrix}^T$$

$$\Phi_C(\mathbf{x}) = \begin{bmatrix} 1 \\ \prod x_i \\ \sqrt{x_1} \end{bmatrix}^T$$

No matter what we choose, the procedure is the same. Replace each \mathbf{x} (row) in our data matrix with $\Phi(\mathbf{x})$, then solve the linear regression problem.

What can we do to reduce overfitting?

- View it as an estimation error
 - *Solution:* Get more data
- Do model selection to find a simpler model
 - Less complicated model, less data needed.
- **Regularization**
 - Add a *prior* belief that the model should be simple!
 - Okay. But how to encode this?





Regularizer View

Add a regularization penalty
 $\lambda \mathbf{w}^T \mathbf{w}$ to the SSE

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{y} - X\mathbf{w})^T (\mathbf{y} - X\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w}$$

Bayesian View

Assume a Gaussian prior
 $\mathbf{w} \sim \mathcal{N}(\vec{0}, \beta I)$


$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmax}} \log P(D|\mathbf{w}) + \log P(\mathbf{w})$$

Arrive at the same solution for L2 regularized least squares.


$$\mathbf{w}^* = (X^T X + \lambda I)^{-1} X^T \mathbf{y}$$

Effect of Regularization - Example 2


Each red line is the result of a 9th degree polynomial fit on 10 random points noisily sampled from the true curve.



$\lambda \approx 13.46$



$\lambda \approx 0.73$



$\lambda \approx 0.09$

Smaller λ resulted in more complex curves that better fit the data but has a high variance.



Regularization is a much more general concept in machine learning. Consider any learning task to minimize a loss $\mathcal{L}(w)$ on the training data - could add a regularizer term.

$$w^* = \underset{w}{\operatorname{argmin}} \mathcal{L}(w) + \lambda * \textit{regularizer}(w)$$

Generally speaking, larger λ 's lead to simpler models and reduce overfitting, smaller λ 's lead to more complex models but improve fit on training data.

Most commonly used regularizers are based on the norm of the weight vector.


Introducing the logistic function:

- May also see it referred to as a sigmoid function or “logit” function

Logistic Function

$$\sigma(z) = \frac{1}{1+e^{-z}}$$

Maps $(-\infty, \infty)$ to $(0,1)$





Logistic Regression Model Assumption:

$$P(y_i = 1 | \mathbf{x}_i; \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

$$P(y_i = 0 | \mathbf{x}_i; \mathbf{w}) = 1 - \sigma(\mathbf{w}^T \mathbf{x}) = \frac{e^{\mathbf{w}^T \mathbf{x}}}{1 + e^{\mathbf{w}^T \mathbf{x}}}$$



Logistic Regression Model's Decision Boundary:

Predict 1 if

$$P(y_i = 1 | \mathbf{x}_i; \mathbf{w}) > P(y_i = 0 | \mathbf{x}_i; \mathbf{w})$$

Divide both sides
by $P(y_i = 0 | \mathbf{x}_i; \mathbf{w})$

$$\Rightarrow \frac{P(y_i = 1 | \mathbf{x}_i; \mathbf{w})}{P(y_i = 0 | \mathbf{x}_i; \mathbf{w})} > 1$$

Just some algebra

$$\Rightarrow e^{\mathbf{w}^T \mathbf{x}} > 1$$


Log of both sides

$$\Rightarrow \mathbf{w}^T \mathbf{x} > 0$$

This tells us the
decision boundary is
the line $\mathbf{w}^T \mathbf{x} = 0$



Example logistic regression decision boundary - linear in the features.



Today's Learning Objectives

Be able to answer:

- What is the perceptron linear classifier?
 - How is it updated?
- How to evaluate classification models?
 - Accuracy
 - Recall and Precision





Perceptron vs Logistic Regression -- History

1830s Logistic function invented


1943 Logistic Regression invented but in biology and largely ignored

1958 Perceptron invented and popularized in AI/ML communities

1970s Logistic Regression reached widespread recognition


Logistic Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Sign Function

$$\text{sign}(z) = \begin{cases} 1 & z \geq 0 \\ -1 & \text{else} \end{cases}$$



Logistic Regression

$$P(y_i = 1 | \mathbf{x}_i; \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}_i) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}_i}}$$

$$y_i = \operatorname{argmax}_y P(y | \mathbf{x}_i; \mathbf{w})$$

Linear Classifier

Probabilistic Interpretation

Trained with gradient descent or other optimization methods

Perceptron

$$y_i = \operatorname{sign}(\mathbf{w}^T \mathbf{x}_i)$$

Linear Classifier

No useful probabilistic interpretation
(mostly just geometry)

Trained with the “perceptron” algorithm



Consider a binary prediction problem where we are given a dataset of \mathbf{x}_i, y_i pairs where $y_i \in \{-1, 1\}$. We want to find optimal linear weights \mathbf{w} such that:

$$y_i = \text{sign}(\mathbf{w}^T \mathbf{x}_i) \quad \forall i$$

This is equivalent to saying we want:


$$y_i \mathbf{w}^T \mathbf{x}_i > 0 \quad \forall i$$



If y_i and $\mathbf{w}^T \mathbf{x}_i$ have same sign, then $\text{sign}(\mathbf{w}^T \mathbf{x}_i)$ will equal y_i



The Perceptron: Setup





Note that the weight vector \mathbf{w} is the *normal vector* of the decision boundary and \mathbf{w} points to the side where $\mathbf{w}^T \mathbf{x} > 0$

Every point where $w^T x = 0$ lies along the line perpendicular to \mathbf{w} . **Why? Linear algebra.**



Linear Algebra Fact: Dot product of two vectors is proportional to the cosine of the angle between them.

$$\vec{v} \cdot \vec{x} = |\vec{x}| |\vec{w}| \cos \theta$$



Trained with a simple iterative algorithm

Perceptron Learning Algorithm

$w = \text{random}(d)$


Repeat:

for each example $x_i, y_i \in D$:

if $y_i w^T x_i < 0$: // if misclassification

$$w = w + \alpha y_i x_i$$

Until no errors or max iterations



The Perceptron: Effect of Update Rule


Consider how adding $\alpha y_i \mathbf{x}_i$ to the weights changes the "error" of a point. Let \mathbf{w}_t be the weights before the update and \mathbf{w}_{t+1} be after the update -- $\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha y_i \mathbf{x}_i$.

$$\begin{aligned} y_i \mathbf{w}_{t+1}^T \mathbf{x}_i &= y_i (\mathbf{w}_t + \alpha y_i \mathbf{x}_i)^T \mathbf{x}_i \\ &= y \mathbf{w}_t^T \mathbf{x}_i + \alpha y_i^2 \mathbf{x}_i^T \mathbf{x}_i \\ &= y \mathbf{w}_t^T \mathbf{x}_i + \alpha y^2 \|\mathbf{x}_i\|_2^2 \end{aligned}$$

This term is a non-negative scalar

$$\Rightarrow y_i \mathbf{w}_{t+1}^T \mathbf{x}_i \geq y_i \mathbf{w}_t^T \mathbf{x}_i$$

We improve the score!



The Perceptron: Effect of Update Rule


Consider how adding $\alpha y_i \mathbf{x}_i$ to the weights changes the “error” of a point. Let \mathbf{w}_t be the weights before the update and \mathbf{w}_{t+1} be after the update -- $\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha y_i \mathbf{x}_i$.

$$\begin{aligned} y_i \mathbf{w}_{t+1}^T \mathbf{x}_i &= y_i (\mathbf{w}_t + \alpha y_i \mathbf{x}_i)^T \mathbf{x}_i \\ &= y \mathbf{w}_t^T \mathbf{x}_i + \alpha y_i^2 \mathbf{x}_i^T \mathbf{x}_i \\ &= y \mathbf{w}_t^T \mathbf{x}_i + \alpha y^2 \|\mathbf{x}_i\|_2^2 \end{aligned}$$

This term is a non-negative scalar

$$\Rightarrow y_i \mathbf{w}_{t+1}^T \mathbf{x}_i \geq y_i \mathbf{w}_t^T \mathbf{x}_i$$

We improve the score!





The Perceptron: Learning Algorithm Revisited

Trained with a simple iterative algorithm

Perceptron Learning Algorithm

$w = \text{random}(d)$


Repeat:

for each example $x_i, y_i \in D$:

if $y_i w^T x_i < 0$: // if misclassification

$$w = w + \alpha y_i x_i$$

Until no errors or max iterations





The Perceptron: Learning Algorithm By Hand


	X1	X2	Y
1	3	4	1
2	6	-3	1
3	-3	9	-1
4	-7	6	-1

$w = [0, 3]$

Repeat until no errors:

for each example $x_i, y_i \in D$:

if $y_i w^T x_i < 0$: $w = w + 1 y_i x_i$





The Perceptron: Learning Algorithm By Hand


	X1	X2	Y
4	3	4	1
2	6	-3	1
1	-3	9	-1
3	-7	6	-1

$w = [-3, 3]$

Repeat until no errors:

for each example $x_i, y_i \in D$:

if $y_i w^T x_i < 0$: $w = w + 1 y_i x_i$





We refer to the perceptron algorithm as an **online learning** algorithm because it updates parameters (aka learns) each time it receives a training example.

In contrast, **batch learning** algorithms collect a set of training examples and learn from them all at once. For example, the gradient descent algorithm we described for logistic regression is a batch algorithm.

(Note that gradient descent also has multiple online versions)



Could modify the perceptron learning algorithm to be a batch learning algorithm by accumulating updates over the whole dataset before changing the weights.

Batched Perceptron Learning Algorithm

\mathbf{w} = random(d)

Repeat:

\mathbf{u} = zero(d)

for each example $\mathbf{x}_i, y_i \in D$:

if $y_i \mathbf{w}^T \mathbf{x}_i < 0$: // if misclassification

$\mathbf{u} += -y_i \mathbf{x}_i$

$\mathbf{w} = \mathbf{w} - \alpha \mathbf{u}$

Until $|\mathbf{u}| \leq \epsilon$

The result looks a lot like gradient descent.

Note for the curious: Perceptron algorithm can be shown to be doing gradient descent to minimize the loss function

$$E(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \max(-y_i \mathbf{w}^T \mathbf{x}_i, 0)$$




Perceptron Learning Algorithm

$w = \text{random}(d)$

While still errors or max iterations:

for each example $x_i, y_i \in D$:

if $y w^T x_i < 0$: $w = w + \alpha y_i x_i$




- Different solutions depending on initialization and order of visiting examples.
- Correcting for a misclassification could move the decision boundary so much that previously correct examples are now misclassified.
 - As such it must go over the training examples multiple times. Each time it goes through the whole training set, it is called an epoch.
- It will terminate if no update is made to w during one epoch



The Perceptron: Limitations

- This algorithm is guaranteed to converge if data is linearly separable – i.e., it reaches a solution in finitely many steps.
 - Why and how many steps? ([proof](#))
- For non-linearly separable cases (like on right), algorithm fails to converge and may be arbitrarily bad if terminated at some maximum number of updates.





The Perceptron: The Voted-Perceptron Algorithm

One Idea: keep around intermediate hypothesis and have them “vote” -- i.e., a weighted average. [Freund and Schapire, 1998]

$w = \text{random}(d)$

$n=0$

While still errors or max iterations:

for each example $x_i, y_i \in D$:

if $y w^T x_i < 0$:

$$w_{n+1} = w_n + \alpha y_i x_i$$

$$c_{n+1} = 0$$

$$n += 1$$

else:

$$c_n += 1$$

Store the weight vector after each error update along with the count of examples that weight was successful before making the error. Weights contribute relative to this value.

$$y_i = \text{sign} \left(\sum_{n=1}^N c_n \text{sign}(w_n^T x_i) \right)$$



- Perceptron incrementally learns a **linear decision** boundary to separate positive from negative (two classes so binary classification).
- It begins with a random or zero weight vector, and incrementally updates the weight vector whenever it makes a mistake.
- For online perceptron, different orderings of the training examples can lead to different outputs. Algorithm will fail to converge if not linearly separable.
- Voted perceptron can handle non-linearly separable data, and is more robust to noise/outlier.

Today's Learning Objectives

Be able to answer:

- ~~What is the perceptron linear classifier?~~
 - ~~How is it updated?~~
- How to evaluate classification models?
 - Accuracy
 - Recall and Precision





We've covered a few predictive models so far in class:

Predictive Models:

Classifiers (given x , produce discrete y)

- k-Nearest Neighbors Classifier
- Logistic Regression (binary classification only)
- Perceptron (binary classification only)

Regressors (given x , produce continuous y)

- Linear Regression
- k-Nearest Neighbors Regressor



Accuracy: $\# \text{ correct} / \text{total}$

If I tell you, I have a ML product that can tell whether someone has a rare disease or not with 99.99% accuracy - are you impressed?

This is a **binary classification** problem - 0 = no disease, 1 = disease

What if the actual percentage of people with the disease is only 0.0001%?

My algorithm is worse than a model that said no-one was a sick which would achieve 99.9999% accuracy.



Let's consider all possible scenarios for a single prediction:

Let y be the true label and \hat{y} be our prediction.

- **False Positive** -- not sick but my model said they have the disease
 $y = 0, \quad \hat{y} = 1$
- **False Negative** -- sick but my model said they don't have the disease
 $y = 1, \quad \hat{y} = 0$
- **True Positive** – sick and my model said they have the disease
 $y = 1, \quad \hat{y} = 1$
- **True Negative** – not sick and model said they don't have the disease
 $y = 0, \quad \hat{y} = 0$



These are often consolidated into what is called a **confusion matrix**.

		Predicted	
		Negative (N) -	Positive (P) +
Actual	Negative -	True Negatives (TN)	False Positives (FP)
	Positive +	False Negatives (FN)	True Positives (TP)



These are often consolidated into what is called a **confusion matrix**.

y	\hat{y}
1	1
0	0
0	1
1	0
1	1
1	1
1	0
1	1

		Predicted	
		Negative (N) -	Positive (P) +
Actual	Negative -	True Negatives (TN)	False Positives (FP)
	Positive +	False Negatives (FN)	True Positives (TP)



Two useful composite measures:

Recall: What fraction of positive does your model predict as positives:

$$\text{Recall} = \frac{\#TP}{\#TP + \#FN}$$

Precision: Of things your model predicts as positives, what fraction are correct?

$$\text{Precision} = \frac{\#TP}{\#TP + \#FP}$$



If I have a model with low recall but high precision, I should trust it _____?

A more when it predicts something as positive than if it predicts it as negative

B the same when it predicts something as positive or negative

C less when it predicts something as positive than if it predicts it as negative

D



Let's go a step further and think about classifiers that output a score:

Let y be the true label and \hat{y} be our prediction. Further let s be the score of our model.

We can write most of our models as doing the following:

$$\hat{y} = \begin{cases} 1 & \text{if } s \geq t \\ 0 & \text{else} \end{cases}$$

For example, linear classifiers like logistic regression or perceptron can be:

$$\hat{y} = \begin{cases} 1 & \text{if } w^T x \geq 0 \\ 0 & \text{else} \end{cases}$$



Evaluating Classifiers

Different settings of the threshold give different outputs.


$$\hat{y} = \begin{cases} 1 & \text{if } s \geq t \\ 0 & \text{else} \end{cases}$$

y	s	$\hat{y} \text{ } t=0$	$\hat{y} \text{ } t=0.25$	$\hat{y} \text{ } t=0.5$	$\hat{y} \text{ } t=0.75$	$\hat{y} \text{ } t=1$
1	0.8	1	1	1	1	0
0	0.2	1	0	0	0	0
0	0.4	1	1	0	0	0
1	0.3	1	1	0	0	0
1	0.7	1	1	1	0	0
1	0.9	1	1	1	1	0
1	0.3	1	1	0	0	0
1	0.5	1	1	1	0	0

And thus different recall/precision values – seems like something we can tune.



Idea: Vary the threshold from its minimum to maximum value and plot the resulting sequence of recall and precision values on a curve.



Call this a precision-vs-recall plot.

Tells us the full range of precision and recall values our classifier can take.

Can be used to compare classifiers or to help us pick optimal settings for a classifier based on our needs.



If a disease can be easily treated early but has devastating effects if given time to develop, we might prefer a classifier with high _____?

A Accuracy

B Precision

C Recall

D



Suppose you have an earthquake warning system with very high recall but low precision. How might people react to this system over time?

Today's Learning Objectives

Be able to answer:

- ~~What is the perceptron linear classifier?~~
 - ~~How is it updated?~~
- ~~How to evaluate classification models?~~
 - ~~Accuracy~~
 - ~~Recall and Precision~~





Next Time: We'll talk about Naïve Bayes and classification problems with more than two classes!